

COMP41450 - Assignment

Kevin Heffernan - ID: 11277955

Introduction For my neural network, I have chosen to use the suggested backpropagation algorithm on the instruction sheet (i.e. for simpler calculation I'm treating the threshold as simply another weight). Results are in bipolar format and I also removed all spaces from the train/test data files which are attached.

Experiment Before adjusting the learning rate, I first needed to find the optimum ranges for the *min* / *max* values used to initialise the *bias* / *weights* for each perceptron. I originally used the same starting ranges for both values but then decided to change them. Since a perceptron cannot fire unless the sum of the weights overcomes the bias ($bias + \sum weights \geq 0$), I set the bias ranges to a much lower value than the weights. I chose $[-1.0 \rightarrow -3.0]$. Next, I needed to find a starting range for the weights. I chose to use $[-0.5 \rightarrow 0.5]$ in proportion to the thresholds. For example, choosing a very large value would result in many perceptrons overcoming the threshold (firing) and resulting in a confusing result. This approach tries to minimise this occurrence but doesn't eliminate the possibility. To address this, I have adopted a "winner takes all" approach in the case where multiple perceptrons have fired. I simply find the perceptrons which fired, and then deactivate all but the one with the highest value > 0 . The next step in the experiment was to find the optimum learning rate. I started at $\alpha = 0.1$ and then moved up incrementally. I also chose to run each test ten times where the correctly classified and incorrectly classified results are taken as averages across the ten tests to account for the fact I'm choosing random values each time from the bias and weight ranges.

Results The data from my search for optimum α values is in figure 1. Looking at these it is clear that choosing a value too low or too high produces a bad result. This is not surprising as a low learning rate makes the network learn very slowly while a high learning rate makes the weights and objective function diverge, so there is no learning at all. Starting out from $\alpha = 0.1$, I then stopped at $\alpha = 0.5$ as this produced a great decrease in accuracy

The best value I found was between $[0.3 \rightarrow 0.4]$. Over my ten tests for each of these values, they both produced a perfect classification. As seen in figure 2, the upper α value produced the greater decrease (steepest slope) so I chose 0.3 as my optimum value.

Conclusion Over the course of my experiment I found that I needed different starting values for my bias and weights. This was done to minimise the occurrence of

Learning Rate	% Correct Class	% Incorrect Class
0.1	76.7	23.3
0.2	83.3	16.7
0.25	86.7	13.3
0.27	96.7	3.3
0.3	100	0
0.4	100	0
0.45	96.7	3.3
0.5	83.3	16.7

Figure 1: results from varied settings of the learning rate




Figure 2: graphical representation of figure 1 results

multiple perceptrons firing. I also found that it was valuable to do ten tests of each setting for the learning rate due to the fact that I was choosing random values at the beginning of each test. Averaging the results then gave me some insight into the true performance of that setting.

Finally, I found that $\alpha = 0.3$ was my optimum value and produced best results. I should also mention that for 10 tests of $\alpha = 0.3$, only one caused a "winner takes all" situation as opposed to 4 instances where $\alpha = 0.5$.